

Camille Harris

Email: camilleharris@berkeley.edu | Website: camille2019.github.io | LinkedIn: linkedin.com/in/camilleharris2019/

Education

Ph.D. Computer Science, Georgia Institute of Technology 2020 - Present.

Advisors: Diyi Yang and Neha Kumar

Thesis Title: An Interdisciplinary Approach to Addressing Discriminatory Outcomes in Downstream Applications of Intelligent Systems

Graduation Date: May 2025

B.A. Computer Science, University of California, Berkeley 2015 - 2019.

Experience

Research Intern at *IBM Research*

May 2023 – August 2023

Tech for Justice Team- Advised by Sara Berger, Raya Horesh, and Rogerio Abreu de Paula. Analysis of the role of data perturbation as a mitigation strategy for gender and ethnicity based-biases in NLP applications.

Research Intern at *Snap Inc.*

May 2022 – Aug 2022

Algorithm Fairness Team, Advised by Aalok Shanbhag and Subhash Sankuratripati. Qualitative research on fairness issues image-processing; Created framework to test Snap camera products across skin tones.

Research Intern at *Adobe Inc.*

May 2020 – Dec 2020

Visualization Service Group, Advised by Ryan Rossi. Created a now patented insight driven visualization recommendation system that created and ranked visualizations for arbitrary dataset.

Select Peer-Reviewed Publications

Exploring Racial and Gender Bias in Automatic Speech Recognition. **Camille Harris**, Chijioke Mgbahurike, Neha Kumar, Diyi Yang. Findings of the Association for Computational Linguistics: Empirical Methods of Natural Language Processing EMNLP 2024

"Honestly, I think TikTok has a Vendetta Against Black Creators": Understanding Black Content Creator

Experiences on TikTok. **Camille Harris**, Amber Johnson, Sadie Palmer, Diyi Yang, Amy Bruckman

The 26th ACM Conference On Computer-Supported Cooperative Work and Social Computing CSCW 2023, Acceptance Rate: 32.7% **Recognition for Diversity and Inclusion**

Exploring the Role of Grammar and Word Choice in Bias Toward African American English (AAE) in Hate speech Classification. **Camille Harris**, Matan Halevy, Ayanna Howard, Amy Bruckman, Diyi Yang. The Fifth Annual ACM Fairness Accountability and Transparency conference. FAccT 2022

Acceptance rate: Overall 36%, Data and Algorithm Evaluation Track 23%

Mitigating Racial Biases in Toxic Language Detection with an Equity-Based Ensemble Framework Matan Halevy, **Camille Harris**, Amy Bruckman, Diyi Yang, Ayanna Howard. ACM Conf. on Equity and Access in Algorithms, Mechanisms, and Optimization. EAAMO 2021

Acceptance rate: 21% **Best Student Paper Award**

Full publication list can be found on Google Scholar.